

Problem Set 10 for lecture Mining Massive Datasets

Due January 20, 2020, 11:59 pm

Exercise 1

(2 points)

Consider a marriage agency in charge of pairing persons based on their similarities and common interests. Each customer has to fill-in b questionnaires Q_1, \dots, Q_b containing r questions each. **Two customers are matched if they give exactly the same answers to all r questions for at least one of these b questionnaires.** Assume that two specific customers C_1 and C_2 give with probability p the same answer to any of the questions in any of the questionnaires. State formulas (and explain how you derived them) for the following cases:

- a) The probability that C_1 and C_2 are matched.
- b) The probability that exactly two (no matter which) questionnaires match, i.e. have the same answers for both C_1 and C_2 .

Exercise 2

(3 points)

Define the graph G_n to have the $2n$ nodes

$$a_0, a_1, \dots, a_{n-1}, b_0, b_1, \dots, b_{n-1}$$

and the following edges. Each a_i (for $i = 0, 1, \dots, n-1$) is connected to the nodes b_j and b_k , where

$$j = 2i \mod n \text{ and } k = (2i + 1) \mod n$$

For instance, the Graph G_4 has the following edges:

$$(a_0, b_0), (a_0, b_1), (a_1, b_2), (a_1, b_3), (a_2, b_0), (a_2, b_1), (a_3, b_2), (a_3, b_3).$$

- a) Find a perfect matching for G_4 and one for G_5 . (1 point)
- b) How many different perfect matchings do G_4 and G_5 have? (2 points)

Exercise 3

(1 point)

Consider the bipartite graph with nodes $1, 2, 3, 4, a, b, c, d$ and edges

$$(1, a), (1, c), (2, b), (3, b), (3, d), (4, a).$$
 Draw this graph.

This bipartite graph has a perfect matching. However, whether or not the greedy algorithm yields a perfect matching depends on the ordering of incoming edges.

Find one ordering of the edges for which the greedy algorithm gives us a perfect matching and one ordering for which it does not.

Exercise 4**(1.5 points)**

Suppose we apply the BALANCE algorithm (with bids of 0 or 1 only) to a situation where advertiser A bids on query words x and y, while advertiser B bids on query words x and z. Both have a budget of \$2. Identify in the following list a sequence of four queries that will certainly be handled optimally by the algorithm, and explain why this is the case (and why not in the other cases).

1) xyyy 2) xyyx 3) yyxx 4) xzyz

Exercise 5**(1.5 points)**

Consider an execution of the BALANCE algorithm with four advertisers, A_1, A_2, A_3, A_4 , and four kinds of queries, Q_1, Q_2, Q_3, Q_4 . Advertiser A_1 bids on queries Q_1 and Q_2 ; A_2 bids on queries Q_2 and Q_3 ; A_3 on queries Q_3 and Q_4 ; and A_4 on queries Q_1 and Q_4 . All bids are equal to 1, and all clickthrough rates are equal. All advertisers have a budget of 3, and ties are broken in favor of the advertiser with the lower index (e.g., A_1 beats A_2). Queries appear in the following order:

$Q_1, Q_2, Q_3, Q_3, Q_1, Q_2, Q_3, Q_1, Q_4, Q_1, Q_4$

What is the sequence of advertisers that the BALANCE algorithm will yield? What is the competitive ratio for this instance?

Exercise 6**(7 points)**

This exercise is the second in the series of tasks related data processing with Spark. You should ideally reuse the implementation developed in the previous problem set (i.e. reading in an *EC2* dataset¹ and creating a Spark dataframe). For each of the following subtasks submit your code as a part of the solution.

- Implement a subroutine in Spark which takes a dataframe (created in the last exercise) as a parameter and generates for each found unique combination ($\langle InstanceType \rangle$, $\langle ProductDescription \rangle$, $\langle AvailabilityZone \rangle$) a new dataframe with associated pairs $\langle Timestamp \rangle$ and $\langle Price \rangle$ (we call such a dataframe a *price timeseries*). Apply Spark's transformations for higher efficiency where possible. The routine should return a dictionary which has as keys combinations ($\langle InstanceType \rangle$, $\langle ProductDescription \rangle$, $\langle AvailabilityZone \rangle$) with each combination encoded as a single string, and as values the corresponding price timeseries.
- Implement a subroutine *saveTimeseries(name, D)* that saves to disk a price timeseries D. The filename should be derived from the string parameter name by respecting the filename restrictions (name is typically a string encoding the combination ($\langle InstanceType \rangle$, $\langle ProductDescription \rangle$, $\langle AvailabilityZone \rangle$) as in part a)). This routine should account for file collisions: if there exist another dataframe D' with the same filename (hence corresponding to the same combination), read this dataframe into memory and merge D with D', removing any duplicates in the process. Save the resulting dataframe into disk.
- To test your implementation, read every file from the provided Amazon dataset and submit the amount of elements and average price obtained for the dataframe with the combination: ($\langle m4.16xlarge \rangle$, $\langle Linux/UNIX \rangle$, $\langle ca-central-1a \rangle$).

¹Available on Moodle, Weekly problem sets: Datasets/dataset-EC2-series/